

Simple Logistic

Andres Potapczynski (ap3635)

```
data <- read_delim(file = file, delim = '|')

## Parsed with column specification:
## cols(
##   .default = col_double(),
##   state = col_character(),
##   county = col_character(),
##   city = col_character(),
##   zip = col_integer(),
##   vendor_name = col_character(),
##   employer_name = col_character(),
##   age = col_integer(),
##   sex = col_character(),
##   sex_F = col_integer(),
##   condition_U = col_integer(),
##   y = col_integer(),
##   y2 = col_integer(),
##   y3 = col_integer(),
##   effective_pay = col_integer(),
##   vendor_Y = col_integer(),
##   employed_30 = col_integer(),
##   antiquity_20 = col_integer(),
##   credit_score = col_integer(),
##   days_wo_pay = col_integer(),
##   months_wo_pay = col_integer()
##   # ... with 12 more columns
## )

## See spec(...) for full column specifications.

# Sample the data
pct = 1
# pct = 0.1
# pct = 0.01
set.seed(seed = 42)
sample_size = round(pct * nrow(data))
sample <- sample(x = nrow(data), size = sample_size, replace = F)
data = data[sample, ]

## Selecting the relevant columns for the analysis
data_sub <- data %>% dplyr::select(
  state,
  city,
  county,
  zip,
  asset_market_value,
  mar_2_app,
  appraisal_value,
  app_2_inc,
  client_income,
```

```

mar_2_inc,
age,
sex_F,
condition_U,
y,
y2)
summary(data_sub)

```

```

##      state           city           county           zip
## Length:30499      Length:30499      Length:30499      Min.   : 1000
## Class :character   Class :character   Class :character   1st Qu.:32680
## Mode  :character   Mode  :character   Mode  :character   Median :55295
##                                           Mean  :54236
##                                           3rd Qu.:76148
##                                           Max.   :99900
## asset_market_value  mar_2_app      appraisal_value      app_2_inc
## Min.   : 120000      Min.   : 0.9184      Min.   : 79521      Min.   :0.004633
## 1st Qu.: 350000      1st Qu.: 1.1259      1st Qu.: 302908      1st Qu.:0.088792
## Median : 398000      Median : 1.2209      Median : 320052      Median :0.100663
## Mean   : 491311      Mean   : 1.3378      Mean   : 371064      Mean   :0.098457
## 3rd Qu.: 463000      3rd Qu.: 1.3620      3rd Qu.: 348289      3rd Qu.:0.111070
## Max.   :4519000      Max.   :21.8468      Max.   :1654602      Max.   :0.367728
## client_income       mar_2_inc       age           sex_F
## Min.   : 143.9      Min.   :0.02738      Min.   :18.00      Min.   :0.0000
## 1st Qu.: 284.3      1st Qu.:0.10934      1st Qu.:27.00      1st Qu.:0.0000
## Median : 310.7      Median :0.12381      Median :32.00      Median :0.0000
## Mean   : 409.8      Mean   :0.12856      Mean   :34.29      Mean   :0.3082
## 3rd Qu.: 341.7      3rd Qu.:0.13978      3rd Qu.:40.00      3rd Qu.:1.0000
## Max.   :1887.2      Max.   :1.09440      Max.   :65.00      Max.   :1.0000
## condition_U         y           y2
## Min.   :0.0000      Min.   :0.00000      Min.   :0.00000
## 1st Qu.:0.0000      1st Qu.:0.00000      1st Qu.:0.00000
## Median :0.0000      Median :0.00000      Median :0.00000
## Mean   :0.3954      Mean   :0.06403      Mean   :0.03371
## 3rd Qu.:1.0000      3rd Qu.:0.00000      3rd Qu.:0.00000
## Max.   :1.0000      Max.   :1.00000      Max.   :1.00000

```

```

## Group data by state and define the IDs
state_summary <- data_sub %>%
  dplyr::select(state,
                client_income,
                appraisal_value,
                asset_market_value) %>%
  group_by(state) %>%
  summarize(n_state = n(),
            income_mean_state = mean(client_income),
            appraisal_mean_state = mean(appraisal_value),
            market_mean_state = mean(asset_market_value)) %>%
  arrange(desc(n_state)) %>%
  ungroup()

```

```

state_summary$ID_state = seq.int(nrow(state_summary))

```

```

## Group data by city and define the IDs

```

```

city_summary <- data_sub %>%
  dplyr::select(city) %>%
  group_by(city) %>%
  summarize(n_city = n()) %>%
  arrange(desc(n_city)) %>%
  ungroup()

city_summary$ID_city = seq.int(nrow(city_summary))

## Merge back into data
data_sub <- data_sub %>%
  inner_join(y = state_summary, by = 'state') %>%
  inner_join(y = city_summary[, c('city', 'ID_city')], by = 'city')

## Rescaling
inputs <- data_sub %>%
  mutate(

    income_st = (log(client_income) - mean(log(client_income))) /
      sd(log(client_income)),

    appraisal_st = (log(appraisal_value) - mean(log(appraisal_value))) /
      sd(log(appraisal_value)),

    market_st = (log(asset_market_value) - mean(log(asset_market_value))) /
      sd(log(asset_market_value)),

    market_state_st = (log(market_mean_state) - mean(log(market_mean_state))) /
      sd(log(market_mean_state)),

    income_state_st = (log(income_mean_state) - mean(log(income_mean_state))) /
      sd(log(income_mean_state)),

    appraisal_state_st = (log(appraisal_mean_state) -
      mean(log(appraisal_mean_state))) /
      sd(log(appraisal_mean_state)),

    mar_2_inc_st = (mar_2_inc - mean(mar_2_inc)) / sd(mar_2_inc),
    app_2_inc_st = (app_2_inc - mean(app_2_inc)) / sd(app_2_inc),
    mar_2_app_st = (mar_2_app - mean(mar_2_app)) / sd(mar_2_app),
    age_st = (age - mean(age)) / sd(age)
  ) %>%
  dplyr::select(
    income_st,
    mar_2_inc_st,
    appraisal_st,
    app_2_inc_st,
    mar_2_app_st,
    market_st,
    age_st,
    income_state_st,
    market_state_st,
    appraisal_state_st,
    y,

```

```

    y2
  )

## Train / Test split
set.seed(seed = 81989843)
pct_train = 0.9
sample_size = round(pct_train * nrow(inputs))
sample <- sample(x = nrow(inputs), size = sample_size, replace = F)

## Allocate train
y = inputs$y

inputs_train = inputs[sample, ]
y_train = y[sample]

## Allocate test
inputs_test = inputs[-sample, ]
y_test = y[-sample]

## Create inputs for STAN
data_stan_train <- list(N=nrow(inputs_train),
                        D=ncol(inputs_train),
                        X=inputs_train,
                        y=y_train)

## Inputs for STAN
X_train = inputs_train %>% dplyr::select(-y, -y2)
X_test = inputs_test %>% dplyr::select(-y, -y2)
N = nrow(X_train)
D = ncol(X_train)
S = length(unique(data_sub$state))
state = data_sub$ID_state[sample]

data_stan_train = list(N=N, D=D, S=S, state=state, X=X_train, y=y_train)

## Inference for Stan model: logistic_base_var_v02.
## 4 chains, each with iter=2000; warmup=1000; thin=1;
## post-warmup draws per chain=1000, total post-warmup draws=4000.
##
##      mean se_mean   sd  2.5%   50%  97.5% n_eff Rhat
## alpha    -2.69    0.00 0.07 -2.82 -2.69 -2.54  1181    1
## beta[1]    0.16    0.01 0.26 -0.35  0.16  0.64  1622    1
## beta[2]    0.20    0.00 0.10  0.01  0.20  0.39  2106    1
## beta[3]    0.55    0.00 0.19  0.17  0.55  0.92  1758    1
## beta[4]   -0.16    0.00 0.11 -0.38 -0.16  0.04  2155    1
## beta[5]    0.10    0.00 0.07 -0.06  0.10  0.22  2174    1
## beta[6]   -0.91    0.00 0.15 -1.20 -0.92 -0.61  3571    1
## beta[7]    0.02    0.00 0.03 -0.04  0.02  0.08  6232    1
## beta[8]    0.11    0.01 0.44 -0.77  0.12  0.95  1206    1
## beta[9]   -0.12    0.01 0.35 -0.79 -0.12  0.55  1463    1
## beta[10]   0.12    0.01 0.49 -0.84  0.11  1.08  1068    1
##
## Samples were drawn using NUTS(diag_e) at Sat Dec  1 10:33:37 2018.
## For each parameter, n_eff is a crude measure of effective sample size,

```

```
## and Rhat is the potential scale reduction factor on split chains (at
## convergence, Rhat=1).
```

```
sims <- rstan::extract(sm.logistic_v01)
alpha_median <- median(sims$alpha)
alpha_s_median <- apply(X = sims$alpha_s, MARGIN = 2, FUN = median)
beta_median <- apply(X = sims$beta, MARGIN = 2, FUN = median)
```

```
threshold = 0.15
```

```
y_hat_baseline <- rep(0, times = length(y_test))
```

```
accuracy_base = sum(y_hat_baseline == y_test) / length(y_test)
```

```
cat('\nBaseline accuracy: ', accuracy_base * 100)
```

```
##
```

```
## Baseline accuracy: 93.54098
```

```
ID_state_test = data_sub$ID_state[-sample]
```

```
proba_hat <- invlogit(as.matrix(X_test) %*% beta_median +
                      alpha_s_median[ID_state_test])
```

```
proba_hat <- as.numeric(proba_hat)
```

```
y_hat = rep(0, times = length(y_test))
```

```
y_hat[proba_hat > threshold] = 1
```

```
accuracy = sum(y_hat == y_test) / length(y_test)
```

```
cat('\nProba max: ', max(proba_hat))
```

```
##
```

```
## Proba max: 0.1814413
```

```
cat('\nLogistic accuracy: ', accuracy * 100)
```

```
##
```

```
## Logistic accuracy: 93.40984
```

```
cat('\nConfusion table\n')
```

```
##
```

```
## Confusion table
```

```
print(table(y_test, y_hat))
```

```
##      y_hat
```

```
## y_test  0    1
```

```
##      0 2848    5
```

```
##      1   196    1
```

```
test_df = data.frame(ID_state = data_sub$ID_state[-sample],
                     state = data_sub$state[-sample],
                     y_test = y_test,
                     y_hat = y_hat)
```

```
test_df <- test_df %>%
  group_by(state) %>%
```

```

summarize(y_sum_test = sum(y_test),
          y_sum_hat = sum(y_hat)) %>%
arrange(desc(y_sum_test)) %>%
ungroup()

accuracy_baseline = mean(abs(test_df$y_sum_test) ** 2)
accuracy = mean(abs(test_df$y_sum_hat - test_df$y_sum_test) ** 2)

cat('\nBaseline MSE: ', accuracy_baseline * 100)

##
## Baseline MSE: 8021.875

cat('\nLogistic MSE: ', accuracy * 100)

##
## Logistic MSE: 8021.875

test_df = data.frame(ID_city = data_sub$ID_city[-sample],
                     city = data_sub$city[-sample],
                     y_test = y_test,
                     y_hat = y_hat)

test_df <- test_df %>%
  group_by(city) %>%
  summarize(y_sum_test = sum(y_test),
            y_sum_hat = sum(y_hat)) %>%
  arrange(desc(y_sum_test)) %>%
  ungroup()

accuracy_baseline = mean(abs(test_df$y_sum_test) ** 2)
accuracy = mean(abs(test_df$y_sum_hat - test_df$y_sum_test) ** 2)

cat('\nBaseline MSE: ', accuracy_baseline * 100)

##
## Baseline MSE: 154.7425

cat('\nLogistic MSE: ', accuracy * 100)

##
## Logistic MSE: 161.2466

y_sum_train = sum(y_train)
y_sum_rep = apply(X = sims$y_rep, MARGIN = 1, FUN = sum)
cat('\nTotal training defaults: ', y_sum_train)

##
## Total training defaults: 1756

cat('\nTotal replicated defaults: ', mean(y_sum_rep))

##
## Total replicated defaults: 1756.493

```