

Zip Code Smoothing

Andres Potapczynski (ap3635)

11/9/2018

Summary

The purpose of this **R**markdown is to see whether it makes sense that the zip codes have different default probabilities or else all can be regularized. The first approach will be to regularized them as if they come from the same distribution. The second approach will be to see if they come from a mixture.

The proposed model is the following

$$y_j \sim \text{Binomial}(n_j, \theta_j)$$

and each θ_j is assumed to come from a Beta distribution.

$$\theta_j \sim \text{Beta}(\alpha, \beta)$$

where the hyperparameters can be obtained via

$$\alpha \sim \text{Ga}(2, 2)$$

and

$$\beta \sim \text{Ga}(2, 2)$$

The second approach has the following generative process [...]

Run the analysis

Load the data

```
file <- './DBs/core.txt'
data <- read_delim(file = file, delim = '|')

# Sample the data
pct = 1
# pct = 0.05
set.seed(seed = 42)
sample_size = round(pct * nrow(data))
sample <- sample(x = nrow(data), size = sample_size, replace = F)
data = data[sample, ]

# Change column format
data$postal_code = factor(data$postal_code)
```

Let's understand the distribution of the zip codes in the DB.

```
zip_summary = data %>%
  # group_by(postal_code, state) %>%
  # group_by(postal_code) %>%
  group_by(postal_code, city, state) %>%
```

```

summarize(mort_no = n(), y_sum = sum(y)) %>%
mutate(theta_emp = y_sum / mort_no) %>%
arrange(desc(mort_no))

```

```

N = zip_summary$mort_no
y = zip_summary$y_sum
M = nrow(zip_summary)

```

The distribution of the zip codes is

```

ggplot(data = zip_summary, mapping = aes(x=theta_emp)) +
  geom_histogram(fill='lightblue', color='black', binwidth = 0.05)

```

```

ggplot(data = zip_summary, mapping = aes(x=postal_code, y=theta_emp)) +
  geom_point()

```

Now, I compile the proposed STAN model.

```

sm <- stan_model('./zip_code_v01.stan')

```

Run first approach

The data for the first model is

```

inputs = list(M=M, N=N, y=y)
model.v01 = sampling(sm, data=inputs)

```

```

## Warning: There were 4 chains where the estimated Bayesian Fraction of Missing Information was low. See
## http://mc-stan.org/misc/warnings.html#bfmi-low

```

```

## Warning: Examine the pairs() plot to diagnose sampling problems

```

```

print(model.v01, digits=2, pars = c('alpha', 'beta'))

```

```

## Inference for Stan model: zip_code_v01.
## 4 chains, each with iter=2000; warmup=1000; thin=1;
## post-warmup draws per chain=1000, total post-warmup draws=4000.
##
##      mean se_mean  sd  2.5%   25%   50%   75%  97.5% n_eff Rhat
## alpha  1.94     0.02 0.16   1.63   1.83   1.93   2.04   2.26    51 1.09
## beta  26.97     0.29 2.34  22.74  25.32  26.82  28.53  31.78    65 1.07
##
## Samples were drawn using NUTS(diag_e) at Fri Nov  9 22:54:33 2018.
## For each parameter, n_eff is a crude measure of effective sample size,
## and Rhat is the potential scale reduction factor on split chains (at
## convergence, Rhat=1).

```

Look at the shrinkage of the model


```

sims = rstan::extract(model.v01)
theta = apply(X = sims$theta, MARGIN = 2, FUN = median)
df = data.frame(theta=theta)
g_post <- ggplot(data = df, mapping = aes(x = theta)) +
  geom_histogram(fill='lightblue', color='black', binwidth = 0.01)
g_post

```

Generate data

```
sims = rstan::extract(model.v01)
y_rep_max = apply(X = sims$y_rep, MARGIN = 1, FUN = max)
df_rep = data.frame(y_rep_max = y_rep_max)
ggplot(data = df_rep, mapping = aes(x = y_rep_max)) +
  geom_histogram(binwidth = 1) +
  geom_vline(xintercept = max(zip_summary$y_sum))
```



Run second approach

[...]

Concluding remarks

[...]