

binomial__baseline

Jongwoo Choi

2018-12-05

```
file <- 'core.txt'
data <- read_delim(file = file, delim = '|')

# Sample the data
pct = 1
# pct = 0.1
# pct = 0.01
set.seed(seed = 42)
sample_size = round(pct * nrow(data))
sample <- sample(x = nrow(data), size = sample_size, replace = F)
data = data[sample, ]
## Selecting the relevant columns for the analysis
data_sub <- data %>% dplyr::select(
  state,
  city,
  county,
  zip,
  asset_market_value,
  mar_2_app,
  appraisal_value,
  app_2_inc,
  client_income,
  mar_2_inc,
  age,
  sex_F,
  condition_U,
  y,
  y2)
summary(data_sub)
```

state	city	county	zip	
Length:30499	Length:30499	Length:30499	Min. : 1000	
Class :character	Class :character	Class :character	1st Qu.:32680	
Mode :character	Mode :character	Mode :character	Median :55295	
			Mean :54236	
			3rd Qu.:76148	
			Max. :99900	
asset_market_value	mar_2_app	appraisal_value	app_2_inc	
Min. : 120000	Min. : 0.9	Min. : 79521	Min. :0.00	
1st Qu.: 350000	1st Qu.: 1.1	1st Qu.: 302908	1st Qu.:0.09	
Median : 398000	Median : 1.2	Median : 320052	Median :0.10	
Mean : 491311	Mean : 1.3	Mean : 371064	Mean :0.10	
3rd Qu.: 463000	3rd Qu.: 1.4	3rd Qu.: 348289	3rd Qu.:0.11	
Max. :4519000	Max. :21.8	Max. :1654602	Max. :0.37	
client_income	mar_2_inc	age	sex_F	condition_U
Min. : 144	Min. :0.03	Min. :18	Min. :0.00	Min. :0.0
1st Qu.: 284	1st Qu.:0.11	1st Qu.:27	1st Qu.:0.00	1st Qu.:0.0

Median :	311	Median :	0.12	Median :	32	Median :	0.00	Median :	0.0
Mean :	410	Mean :	0.13	Mean :	34	Mean :	0.31	Mean :	0.4
3rd Qu.:	342	3rd Qu.:	0.14	3rd Qu.:	40	3rd Qu.:	1.00	3rd Qu.:	1.0
Max. :	1887	Max. :	1.09	Max. :	65	Max. :	1.00	Max. :	1.0

	y	y2
Min. :	0.00	Min. :0.00
1st Qu.:	0.00	1st Qu.:0.00
Median :	0.00	Median :0.00
Mean :	0.06	Mean :0.03
3rd Qu.:	0.00	3rd Qu.:0.00
Max. :	1.00	Max. :1.00

```
## Group data by state and define the IDs
```

```
state_summary <- data_sub %>%
  dplyr::select(state,
                client_income,
                appraisal_value,
                asset_market_value) %>%
  group_by(state) %>%
  summarize(n_state = n(),
            income_mean_state = mean(client_income),
            appraisal_mean_state = mean(appraisal_value),
            market_mean_state = mean(asset_market_value)) %>%
  arrange(desc(n_state)) %>%
  ungroup()
state_summary$ID_state = seq.int(nrow(state_summary))
```

```
## Group data by city and define the IDs
```

```
city_summary <- data_sub %>%
  dplyr::select(city, state,
                client_income,
                appraisal_value,
                asset_market_value,
                mar_2_inc,
                mar_2_app,
                app_2_inc,
                age,
                y,
                y2) %>%
  group_by(city, state) %>%
  summarize(n_city = n(),
            income_mean_city = mean(client_income),
            appraisal_mean_city = mean(appraisal_value),
            market_mean_city = mean(asset_market_value),
            mar_2_inc_mean_city = mean(mar_2_inc),
            mar_2_app_mean_city = mean(mar_2_app),
            app_2_inc_mean_city = mean(app_2_inc),
            age_mean_city = mean(age),
            sum_y = sum(y),
            sum_y2 = sum(y2)) %>%
  arrange(desc(n_city)) %>%
  ungroup()
```

```

## Merge back into data
city_summary <- city_summary %>%
  inner_join(y = state_summary[c('ID_state', 'state')], by = 'state')

## Rescaling
inputs <- city_summary %>%
  mutate(
    market_state_city = (log(market_mean_city) - mean(log(market_mean_city))) /
      sd(log(market_mean_city)),

    income_state_city = (log(appraisal_mean_city) - mean(log(appraisal_mean_city))) /
      sd(log(appraisal_mean_city)),

    appraisal_state_city = (log(appraisal_mean_city) -
      mean(log(appraisal_mean_city))) /
      sd(log(appraisal_mean_city)),

    mar_2_inc_city = (mar_2_inc_mean_city - mean(mar_2_inc_mean_city)) / sd(mar_2_inc_mean_city),

    app_2_inc_city = (app_2_inc_mean_city - mean(app_2_inc_mean_city)) / sd(app_2_inc_mean_city),

    mar_2_app_city = (mar_2_app_mean_city - mean(mar_2_app_mean_city)) / sd(mar_2_app_mean_city),

    age_city = (age_mean_city - mean(age_mean_city)) / sd(age_mean_city)) %>%
  dplyr::select(
    market_state_city,
    income_state_city,
    appraisal_state_city,
    mar_2_inc_city,
    app_2_inc_city,
    mar_2_app_city,
    age_city,
    ID_state,
    n_city,
    sum_y,
    sum_y2
  )

```

Baseline Model: Binomial

Our baseline model is simple binomial regression model. We give weak cauchy priors on the coefficient parameter β the intercept a . In the city level, we assume that the number of individual records in city i is n_i .

$$a \sim \text{Cauchy}(0, 10)$$

$$\beta \sim \text{Cauchy}(0, 2.5)$$

$$y_i \sim \text{Binomial}(n_i, \text{logit}^{-1}(a + \beta \cdot X_i))$$

The binomial baseline model is given by:

```

baseline_model_binom=stan_model('binomial_baseline_city.stan')

print_file('binomial_baseline_city.stan')

// baseline model: city level
data {
  int<lower=1> N_train;           // number of record, train
  int<lower=1> N_test;           // number of record, test
  int<lower=1> D;                // number of covariates
  matrix[N_train, D] X_train;   // train data
  matrix[N_test, D] X_test;     // test data
  int<lower=1> n_city_train[N_train]; // number of record for city n, train
  int<lower=1> n_city_test[N_test]; // number of record for city n, test
  int<lower=0> y_train[N_train]; // y train
}
parameters {
  real a;                       // include intercept
  vector[D] beta;               // regression coefficient vector
}
transformed parameters {
  vector[N_train] eta;
  eta = a + X_train*beta;       // probability in binomial regression
}
model {
  a ~ cauchy(0, 10);            // Cauchy prior
  beta ~ cauchy(0, 2.5);        // Cauchy prior
  y_train ~ binomial_logit(n_city_train, eta); // binomial model
}
generated quantities{
  int<lower =0> y_rep[N_train];
  int<lower =0> y_rep_cv[N_test];
  for (i in 1:N_train){
    y_rep[i] = binomial_rng(n_city_train[i], inv_logit(eta[i]));
  }
  for (i in 1:N_test){
    y_rep_cv[i] = binomial_rng(n_city_test[i], inv_logit(eta[i]));
  }
}

```

Parameters recovered

We generate the fake data to simulate the model.

```

a <- rcauchy(1, 0, 10)
beta <- rcauchy(7, 0, 2.5)
X <- inputs %>% dplyr::select(-ID_state, -n_city, -sum_y, -sum_y2)
n_city <- inputs$n_city
N = nrow(X)
D = ncol(X)

y_fake <- c()
for (i in 1:N){
  y_fake[i] <- rbinom(1, n_city[i],
                     invlogit(a + beta %*% as.matrix(X)[i,]))
}

```

```

}

fake_baseline_data <- list(N_train=N, N_test=N, D=D,
                           X_train=X, X_test=X,
                           n_city_train = n_city,
                           n_city_test = n_city,
                           y_train = y_fake)

fit_fake <- sampling(baseline_model_binom,
                    data=fake_baseline_data, seed=1234)
saveRDS(fit_fake, file = 'baseline_fit_fake.rds')

fit_fake <- readRDS(file = 'baseline_fit_fake.rds')
print(fit_fake, pars = c('a', 'beta'),
      digits = 2, probs = c(0.025, 0.5, 0.975))


```

Inference for Stan model: binomial_baseline_city.
 4 chains, each with iter=2000; warmup=1000; thin=1;
 post-warmup draws per chain=1000, total post-warmup draws=4000.

	mean	se_mean	sd	2.5%	50%	98%	n_eff	Rhat
a	-9.63	0.00	0.19	-9.99	-9.63	-9.3	2317	1
beta[1]	4.53	0.02	0.90	2.79	4.52	6.3	2545	1
beta[2]	-0.33	0.09	2.55	-5.32	-0.37	4.7	837	1
beta[3]	-0.33	0.09	2.53	-5.45	-0.29	4.5	846	1
beta[4]	1.40	0.02	0.78	-0.13	1.39	2.9	1831	1
beta[5]	1.15	0.01	0.53	0.13	1.15	2.2	1802	1
beta[6]	-23.05	0.02	0.94	-24.94	-23.04	-21.2	1496	1
beta[7]	0.68	0.00	0.11	0.45	0.68	0.9	4000	1

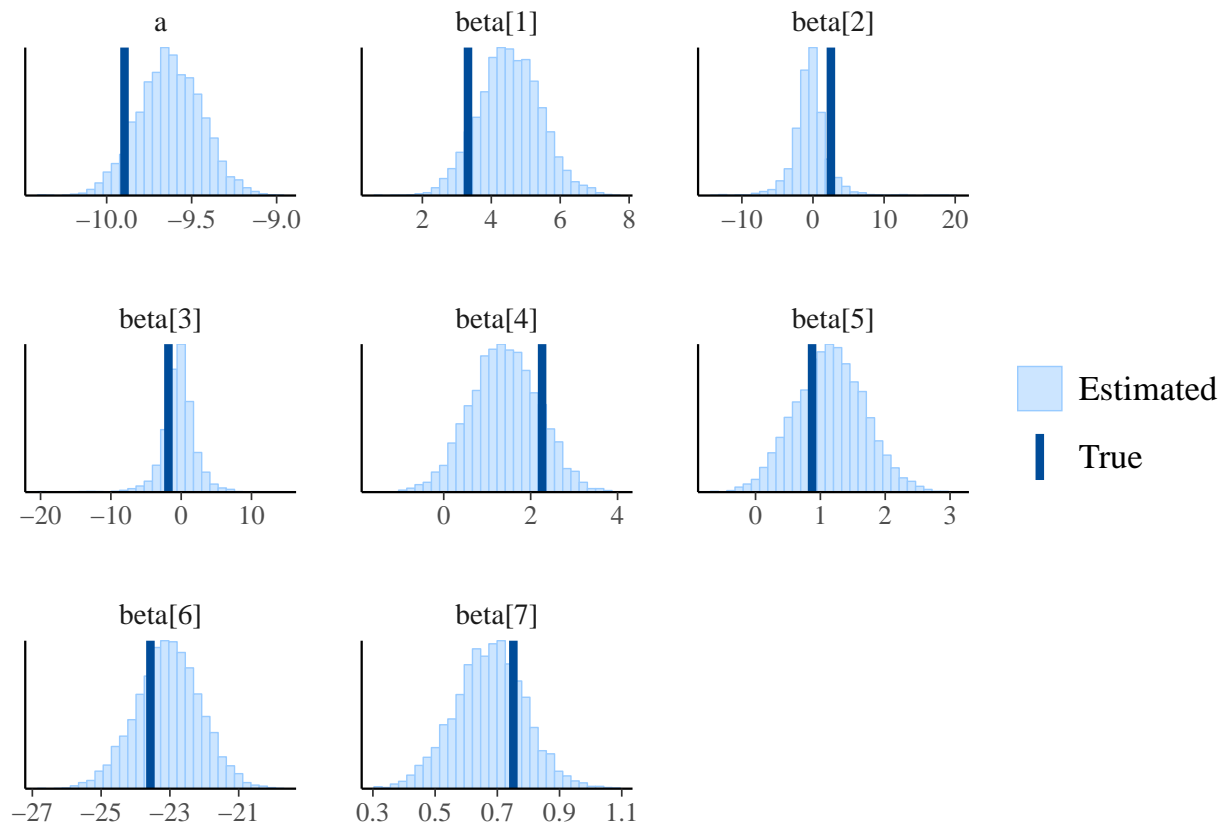
Samples were drawn using NUTS(diag_e) at Wed Dec 5 20:20:24 2018.
 For each parameter, n_eff is a crude measure of effective sample size,
 and Rhat is the potential scale reduction factor on split chains (at
 convergence, Rhat=1).

```
plot(fit_fake)
```



The following plot shows how well parameters recovered.

```
sims_fake <- as.matrix(fit_fake)
true <- c(a, beta)
color_scheme_set("brightblue")
mcmc_recover_hist(sims_fake[, 1:8], true)
```



We see that most β 's didn't recovered well except β_6 and β_7 .

Fit real data

Now we fit the real data with this model.

```
y = inputs$sum_y
#y2 = inputs$sum_y2

baseline_data = list(N_train=N, N_test=N, D=D,
                     X_train=X, X_test=X,
                     n_city_train = n_city,
                     n_city_test = n_city,
                     y_train = y)

fit1 <- sampling(baseline_model_binom,
                 data=baseline_data, seed=1234)
saveRDS(fit1, file = 'baseline_fit1.rds')
```

The following shows the posterior mean and sd of our parameters.

```
fit1 <- readRDS(file = 'baseline_fit1.rds')
print(fit1, pars=c('a', 'beta', 'lp__'),
      digits = 2, probs = c(0.025, 0.5, 0.975))
```

Inference for Stan model: binomial_baseline_city.
 4 chains, each with iter=2000; warmup=1000; thin=1;
 post-warmup draws per chain=1000, total post-warmup draws=4000.

	mean	se_mean	sd	2.5%	50%	98%	n_eff	Rhat
a	-2.68	0.00	0.04	-2.76	-2.68	-2.61	2891	1
beta[1]	-0.69	0.01	0.36	-1.39	-0.68	0.01	2346	1
beta[2]	0.45	0.05	2.34	-4.45	0.33	5.47	1859	1
beta[3]	0.29	0.05	2.33	-4.67	0.39	5.13	1844	1
beta[4]	0.47	0.01	0.22	0.03	0.48	0.90	1985	1
beta[5]	-0.44	0.00	0.18	-0.78	-0.44	-0.09	1996	1
beta[6]	-0.14	0.00	0.20	-0.55	-0.14	0.24	2305	1
beta[7]	-0.08	0.00	0.08	-0.24	-0.08	0.08	3162	1
lp__	-7252.01	0.06	2.11	-7256.99	-7251.65	-7248.94	1267	1

Samples were drawn using NUTS(diag_e) at Wed Dec 5 20:26:39 2018.

For each parameter, `n_eff` is a crude measure of effective sample size, and `Rhat` is the potential scale reduction factor on split chains (at convergence, `Rhat`=1).

The posterior mean of intercept a is -2.68 with standard deviation 0.04 . We can also see the posterior mean and standard deviation of coefficient parameter β from the above table. We see that β_2 and β_3 have larger standard deviation compare to others.

PPC

In stan model, we generated y_{rep} of city level using `binomial_rng` function in stan.


```
'transformed parameters {
  vector[N_train] eta;
  eta = a + X_train*beta;          // probability in binomial regression
}

generated quantities{
  int<lower =0> y_rep[N_train];
  for (i in 1:N_train){
    y_rep[i] = binomial_rng(n_city_train[i], inv_logit(eta[i]));
  }
}'
```

```
[1] "transformed parameters {\n  vector[N_train] eta;\n  eta = a + X_train*beta;          // probab"
```

The following shows the posterior predictive check.

```
y_rep <- as.matrix(fit1, pars = "y_rep")
ppc_dens_overlay(y = y, y_rep[1:200,])
```

We see that the new y fit well even though there were some gaps.

```
sims <- rstan::extract(fit1)

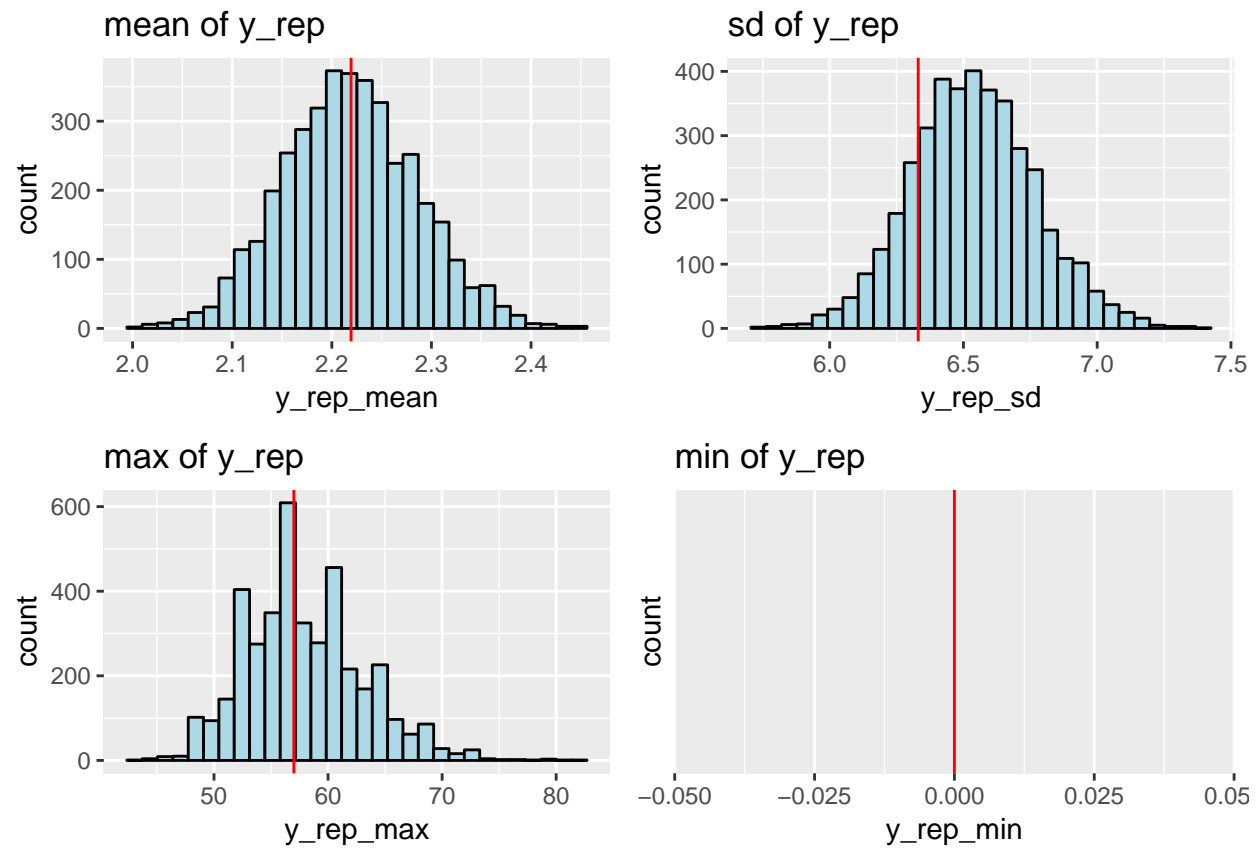
df <- data.frame(y_rep_mean = apply(X=sims$y_rep, MARGIN = 1, FUN = mean))
meangg <- ggplot(df, aes(x=y_rep_mean)) +
  geom_histogram(fill='lightblue', color='black') +
  geom_vline(xintercept = mean(y), color='red') +
  ggtitle('mean of y_rep')

df <- data.frame(y_rep_sd = apply(X = sims$y_rep, MARGIN = 1, FUN = sd))
sdgg <- ggplot(df, aes(x=y_rep_sd)) +
  geom_histogram(fill='lightblue', color='black') +
  geom_vline(xintercept = sd(y), color='red') +
  ggtitle('sd of y_rep')


df <- data.frame(y_rep_max = apply(X = sims$y_rep, MARGIN = 1, FUN = max))
maxgg <- ggplot(df, aes(x=y_rep_max)) +
  geom_histogram(fill='lightblue',color='black') +
  geom_vline(xintercept = max(y), color='red') +
  ggtitle('max of y_rep')

df <- data.frame(y_rep_min = apply(X = sims$y_rep, MARGIN = 1, FUN = min))
mingg <- ggplot(df, aes(x=y_rep_min)) +
  geom_histogram(fill='lightblue',color='black') +
  geom_vline(xintercept = min(y), color='red') +
  ggtitle('min of y_rep')
```

```
gridExtra::grid.arrange(meangg, sdgg, maxgg, mingg,
  layout_matrix = rbind(c(1, 2),
    c(3, 4)))
```



```
# Scatterplot of two test statistics
ppc_scatter_avg(y = y, yrep = y_rep)
```



Scatter plot looks linear but it is not perfect.

Evaluation (RMSE)

We evaluate our training model with RMSE using 5 fold cross validation.

```
## K fold CV
set.seed(1234)
splited_inputs <- split(inputs, sample(rep(1:5, 176)))

for (i in 1:5){
  a <- c(1,2,3,4,5)[-i]
  inputs_test = splited_inputs[[i]]
  inputs_train = rbind(splited_inputs[[a[1]]],
                        splited_inputs[[a[2]]],
                        splited_inputs[[a[3]]],
                        splited_inputs[[a[4]]])

  y_train = inputs_train$sum_y
  y_test = inputs_test$sum_y

  ## Inputs for STAN
  n_city_train = inputs_train$n_city
  n_city_test = inputs_test$n_city

  X_train = inputs_train %>%
```

```

dplyr::select(-ID_state,-n_city,-sum_y, -sum_y2)
X_test = inputs_test %>%
  dplyr::select(-ID_state,-n_city,-sum_y, -sum_y2)

N_train = nrow(X_train)
N_test = nrow(X_test)

D = ncol(X_train)

baseline_data = list(N_train=N_train, N_test=N_test, D=D,
                      X_train=X_train, X_test=X_test,
                      n_city_train = n_city_train,
                      n_city_test = n_city_test,
                      y_train = y_train)
fit_cv <- sampling(baseline_model_binom, data=baseline_data, seed=1234)
name = paste('model1_cv', as.character(i), '.rds', sep = "")
saveRDS(fit_cv, file = name)
}

```

```

rmse <- c()
for (i in 1:5){
  name = paste('model1_cv', as.character(i), '.rds', sep = "")
  fit_cv <- readRDS(file = name)
  sims_cv <- rstan::extract(fit_cv)
  y_hat <- apply(X = sims_cv$y_rep_cv, MARGIN = 2, FUN = median)

  test_df = data.frame(ID_state = inputs_test$ID_state,
                       y_test = y_test,
                       y_hat = y_hat)

  test_df <- test_df %>%
    summarize(y_sum_test = sum(y_test),
              y_sum_hat = sum(y_hat)) %>%
    arrange(desc(y_sum_test)) %>%
    ungroup()

  #mse_baseline = mean((test_df$y_sum_test) ** 2)
  rmse[i] = sqrt(mean((test_df$y_sum_hat - test_df$y_sum_test) ** 2))
}

average_RMSE <- mean(rmse)
sd_RMSE <- sd(rmse)

cat('The average of RMSE of baseline model is: ', average_RMSE)

```

The average of RMSE of baseline model is: 75

```
cat('The standard deviation of RMSE of baseline model is: ', sd_RMSE)
```

The standard deviation of RMSE of baseline model is: 34