

binomial__baseline

Jongwoo Choi

2018-12-03

```
file <- 'core.txt'
data <- read_delim(file = file, delim = '|')

# Sample the data
pct = 1
# pct = 0.1
# pct = 0.01
set.seed(seed = 42)
sample_size = round(pct * nrow(data))
sample <- sample(x = nrow(data), size = sample_size, replace = F)
data = data[sample, ]
## Selecting the relevant columns for the analysis
data_sub <- data %>% dplyr::select(
  state,
  city,
  county,
  zip,
  asset_market_value,
  mar_2_app,
  appraisal_value,
  app_2_inc,
  client_income,
  mar_2_inc,
  age,
  sex_F,
  condition_U,
  y,
  y2)
summary(data_sub)
```

| state | city | county | zip | |
|--------------------|------------------|------------------|---------------|-------------|
| Length:30499 | Length:30499 | Length:30499 | Min. : 1000 | |
| Class :character | Class :character | Class :character | 1st Qu.:32680 | |
| Mode :character | Mode :character | Mode :character | Median :55295 | |
| | | | Mean :54236 | |
| | | | 3rd Qu.:76148 | |
| | | | Max. :99900 | |
| asset_market_value | mar_2_app | appraisal_value | app_2_inc | |
| Min. : 120000 | Min. : 0.9 | Min. : 79521 | Min. :0.00 | |
| 1st Qu.: 350000 | 1st Qu.: 1.1 | 1st Qu.: 302908 | 1st Qu.:0.09 | |
| Median : 398000 | Median : 1.2 | Median : 320052 | Median :0.10 | |
| Mean : 491311 | Mean : 1.3 | Mean : 371064 | Mean :0.10 | |
| 3rd Qu.: 463000 | 3rd Qu.: 1.4 | 3rd Qu.: 348289 | 3rd Qu.:0.11 | |
| Max. :4519000 | Max. :21.8 | Max. :1654602 | Max. :0.37 | |
| client_income | mar_2_inc | age | sex_F | condition_U |
| Min. : 144 | Min. :0.03 | Min. :18 | Min. :0.00 | Min. :0.0 |
| 1st Qu.: 284 | 1st Qu.:0.11 | 1st Qu.:27 | 1st Qu.:0.00 | 1st Qu.:0.0 |

| | | | | | | | | | |
|----------|------|----------|------|----------|----|----------|------|----------|-----|
| Median : | 311 | Median : | 0.12 | Median : | 32 | Median : | 0.00 | Median : | 0.0 |
| Mean : | 410 | Mean : | 0.13 | Mean : | 34 | Mean : | 0.31 | Mean : | 0.4 |
| 3rd Qu.: | 342 | 3rd Qu.: | 0.14 | 3rd Qu.: | 40 | 3rd Qu.: | 1.00 | 3rd Qu.: | 1.0 |
| Max. : | 1887 | Max. : | 1.09 | Max. : | 65 | Max. : | 1.00 | Max. : | 1.0 |

| | y | y2 |
|----------|------|--------------|
| Min. : | 0.00 | Min. :0.00 |
| 1st Qu.: | 0.00 | 1st Qu.:0.00 |
| Median : | 0.00 | Median :0.00 |
| Mean : | 0.06 | Mean :0.03 |
| 3rd Qu.: | 0.00 | 3rd Qu.:0.00 |
| Max. : | 1.00 | Max. :1.00 |

```
## Group data by state and define the IDs
state_summary <- data_sub %>%
  dplyr::select(state,
                client_income,
                appraisal_value,
                asset_market_value) %>%
  group_by(state) %>%
  summarize(n_state = n(),
            income_mean_state = mean(client_income),
            appraisal_mean_state = mean(appraisal_value),
            market_mean_state = mean(asset_market_value)) %>%
  arrange(desc(n_state)) %>%
  ungroup()
state_summary$ID_state = seq.int(nrow(state_summary))
```

```
## Group data by city and define the IDs
city_summary <- data_sub %>%
  dplyr::select(city, state,
                client_income,
                appraisal_value,
                asset_market_value,
                mar_2_inc,
                mar_2_app,
                app_2_inc,
                age,
                y,
                y2) %>%
  group_by(city, state) %>%
  summarize(n_city = n(),
            income_mean_city = mean(client_income),
            appraisal_mean_city = mean(appraisal_value),
            market_mean_city = mean(asset_market_value),
            mar_2_inc_mean_city = mean(mar_2_inc),
            mar_2_app_mean_city = mean(mar_2_app),
            app_2_inc_mean_city = mean(app_2_inc),
            age_mean_city = mean(age),
            sum_y = sum(y),
            sum_y2 = sum(y2)) %>%
  arrange(desc(n_city)) %>%
  ungroup()
```

```

## Merge back into data
city_summary <- city_summary %>%
  inner_join(y = state_summary[c('ID_state', 'state')], by = 'state')

## Rescaling
inputs <- city_summary %>%
  mutate(
    market_state_city = (log(market_mean_city) - mean(log(market_mean_city))) /
      sd(log(market_mean_city)),

    income_state_city = (log(appraisal_mean_city) - mean(log(appraisal_mean_city))) /
      sd(log(appraisal_mean_city)),

    appraisal_state_city = (log(appraisal_mean_city) -
      mean(log(appraisal_mean_city))) /
      sd(log(appraisal_mean_city)),

    mar_2_inc_city = (mar_2_inc_mean_city - mean(mar_2_inc_mean_city)) / sd(mar_2_inc_mean_city),

    app_2_inc_city = (app_2_inc_mean_city - mean(app_2_inc_mean_city)) / sd(app_2_inc_mean_city),

    mar_2_app_city = (mar_2_app_mean_city - mean(mar_2_app_mean_city)) / sd(mar_2_app_mean_city),

    age_city = (age_mean_city - mean(age_mean_city)) / sd(age_mean_city)) %>%
  dplyr::select(
    market_state_city,
    income_state_city,
    appraisal_state_city,
    mar_2_inc_city,
    app_2_inc_city,
    mar_2_app_city,
    age_city,
    ID_state,
    n_city,
    sum_y,
    sum_y2
  )

```

Baseline Model: Binomial

Our baseline model is binomial. We give cauchy priors on the coefficient parameter β the intercept a . In the city level, we assume that the number of individual records in city i is n_i .

$$a \sim \text{Cauchy}(0, 10)$$

$$\beta \sim \text{Cauchy}(0, 2.5)$$

$$y_i \sim \text{Binomial}(n_i, \text{logit}^{-1}(a + \beta \cdot X_i))$$

The binomial baseline model is given by:

```

print_file('binomial_baseline_city.stan')

// baseline model: city level
data {
  int<lower=1> N_train;           // number of record, train
  int<lower=1> N_test;           // number of record, test
  int<lower=1> D;                // number of covariates
  matrix[N_train, D] X_train;   // train data
  matrix[N_test, D] X_test;     // test data
  int<lower=1> n_city_train[N_train]; // number of record for city n, train
  int<lower=1> n_city_test[N_test];  // number of record for city n, test
  int<lower=0> y_train[N_train];    // y train
}
parameters {
  // regression coefficient vector
  real a; // include intercept
  vector[D] beta;
}
transformed parameters {
  vector[N_train] eta;
  eta = a + X_train*beta;
}
model {
  a ~ normal(0, 5);
  beta ~ normal(0, 5);
  y_train ~ binomial_logit(n_city_train, eta);
}
generated quantities{
  int<lower =0> y_rep[N_train];
  int<lower =0> y_rep_cv[N_test];
  for (i in 1:N_train){
    y_rep[i] = binomial_rng(n_city_train[i], inv_logit(eta[i]));
  }
  for (i in 1:N_test){
    y_rep_cv[i] = binomial_rng(n_city_test[i], inv_logit(eta[i]));
  }
}

```

Parameters recovered

We generate the fake data to simulate the model.

```

a <- rcauchy(1, 0, 10)
beta <- rcauchy(7, 0, 2.5)
X <- inputs %>% dplyr::select(-ID_state, -n_city, -sum_y, -sum_y2)
n_city <- inputs$n_city
N = nrow(X)
D = ncol(X)

y_fake <- c()
for (i in 1:N){
  y_fake[i] <- rbinom(1, n_city[i],
                     invlogit(a + beta %*% as.matrix(X)[i,]))
}

```

```

}

fake_baseline_data <- list(N_train=N, N_test=N, D=D,
                           X_train=X, X_test=X,
                           n_city_train = n_city,
                           n_city_test = n_city,
                           y_train = y_fake)

baseline_model_binom=stan_model('binomial_baseline_city.stan')
fit_fake <- sampling(baseline_model_binom,
                    data=fake_baseline_data, seed=1234)

print(fit_fake, pars = c('a', 'beta'),
      digits = 2, probs = c(0.025, 0.5, 0.975))


```

Inference for Stan model: binomial_baseline_city.
 4 chains, each with iter=2000; warmup=1000; thin=1;
 post-warmup draws per chain=1000, total post-warmup draws=4000.

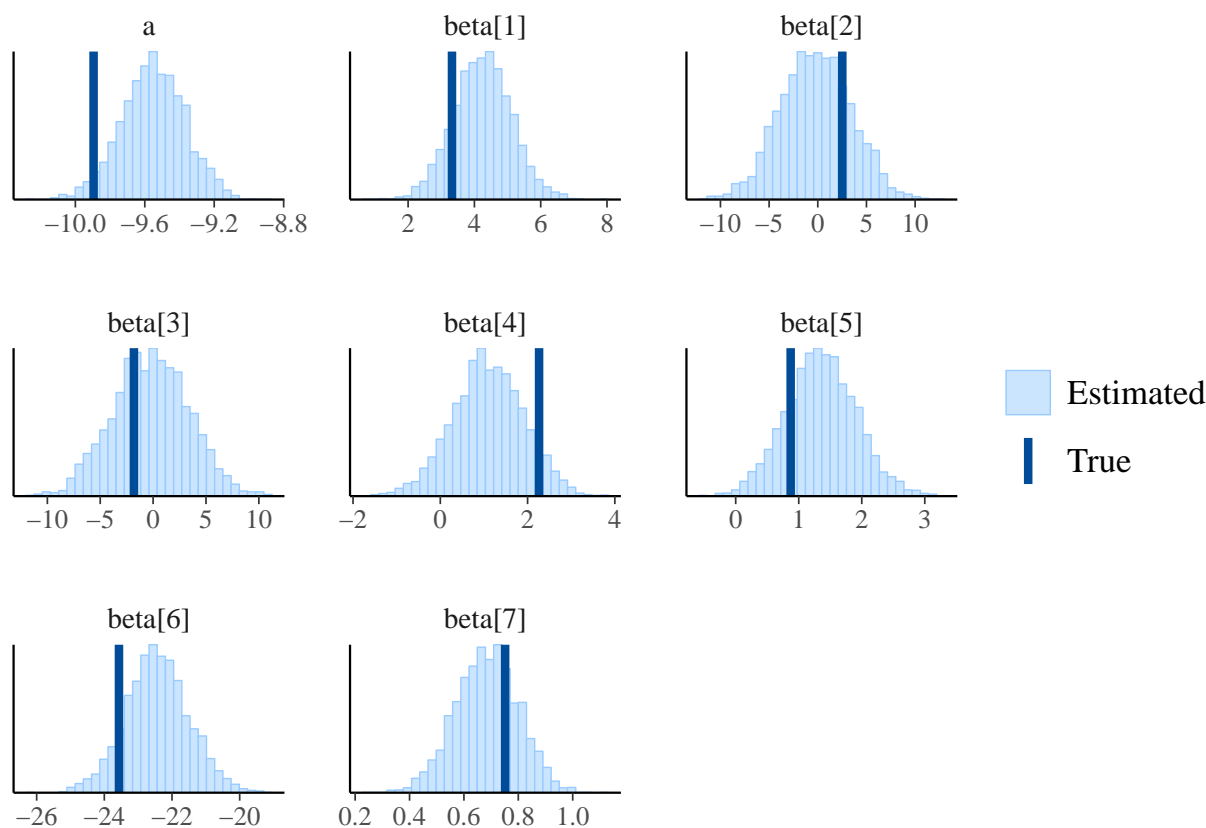
| | mean | se_mean | sd | 2.5% | 50% | 98% | n_eff | Rhat |
|---------|--------|---------|------|--------|--------|--------|-------|------|
| a | -9.55 | 0.00 | 0.18 | -9.92 | -9.55 | -9.19 | 3065 | 1 |
| beta[1] | 4.24 | 0.02 | 0.88 | 2.51 | 4.25 | 5.99 | 2739 | 1 |
| beta[2] | -0.21 | 0.07 | 3.63 | -7.30 | -0.21 | 6.68 | 2367 | 1 |
| beta[3] | -0.21 | 0.07 | 3.59 | -7.15 | -0.21 | 6.78 | 2388 | 1 |
| beta[4] | 1.09 | 0.02 | 0.82 | -0.54 | 1.09 | 2.67 | 1893 | 1 |
| beta[5] | 1.35 | 0.01 | 0.56 | 0.26 | 1.35 | 2.48 | 1898 | 1 |
| beta[6] | -22.43 | 0.02 | 0.94 | -24.32 | -22.45 | -20.55 | 1981 | 1 |
| beta[7] | 0.68 | 0.00 | 0.12 | 0.46 | 0.68 | 0.91 | 3455 | 1 |

Samples were drawn using NUTS(diag_e) at Mon Dec 3 19:30:32 2018.
 For each parameter, n_eff is a crude measure of effective sample size,
 and Rhat is the potential scale reduction factor on split chains (at
 convergence, Rhat=1).

```
plot(fit_fake)
```



```
sims_fake <- as.matrix(fit_fake)
true <- c(a, beta)
color_scheme_set("brightblue")
mcmc_recover_hist(sims_fake[, 1:8], true)
```



Fit real data

Now we fit the real data with this model.

```
y = inputs$sum_y
#y2 = inputs$sum_y2

baseline_data = list(N_train=N, N_test=N, D=D,
                     X_train=X, X_test=X,
                     n_city_train = n_city,
                     n_city_test = n_city,
                     y_train = y)
```

```
fit1 <- sampling(baseline_model_binom,
                 data=baseline_data, seed=1234)
```

```
print(fit1, pars=c('a', 'beta', 'lp__'),
      digits = 2, probs = c(0.025, 0.5, 0.975))
```

Inference for Stan model: binomial_baseline_city.

4 chains, each with iter=2000; warmup=1000; thin=1;

post-warmup draws per chain=1000, total post-warmup draws=4000.

| | mean | se_mean | sd | 2.5% | 50% | 98% | n_eff | Rhat |
|---------|-------|---------|------|-------|-------|-------|-------|------|
| a | -2.68 | 0.00 | 0.04 | -2.76 | -2.68 | -2.61 | 2547 | 1 |
| beta[1] | -0.73 | 0.01 | 0.38 | -1.47 | -0.73 | 0.01 | 2140 | 1 |
| beta[2] | 0.31 | 0.08 | 3.67 | -6.89 | 0.38 | 7.68 | 1876 | 1 |
| beta[3] | 0.46 | 0.08 | 3.65 | -6.80 | 0.44 | 7.82 | 1902 | 1 |


| | | | | | | | | |
|---------|----------|------|------|----------|----------|----------|------|---|
| beta[4] | 0.49 | 0.01 | 0.23 | 0.04 | 0.48 | 0.95 | 1789 | 1 |
| beta[5] | -0.45 | 0.00 | 0.18 | -0.82 | -0.45 | -0.09 | 1719 | 1 |
| beta[6] | -0.14 | 0.00 | 0.21 | -0.57 | -0.13 | 0.27 | 2094 | 1 |
| beta[7] | -0.08 | 0.00 | 0.08 | -0.24 | -0.08 | 0.08 | 2811 | 1 |
| lp__ | -7251.78 | 0.05 | 2.03 | -7256.58 | -7251.42 | -7248.89 | 1536 | 1 |

Samples were drawn using NUTS(diag_e) at Mon Dec 3 19:34:02 2018.

For each parameter, n_eff is a crude measure of effective sample size, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat=1).

PPC

```
y_rep <- as.matrix(fit1, pars = "y_rep")
ppc_dens_overlay(y = y, y_rep[1:200,])
```



```
sims <- rstan::extract(fit1)

df <- data.frame(y_rep_mean = apply(X=sims$y_rep, MARGIN = 1, FUN = mean))
meangg <- ggplot(df, aes(x=y_rep_mean)) +
  geom_histogram(fill='lightblue', color='black') +
  geom_vline(xintercept = mean(y), color='red') +
  ggtitle('mean of y_rep')

df <- data.frame(y_rep_sd = apply(X = sims$y_rep, MARGIN = 1, FUN = sd))
sdgg <- ggplot(df, aes(x=y_rep_sd)) +
```



```

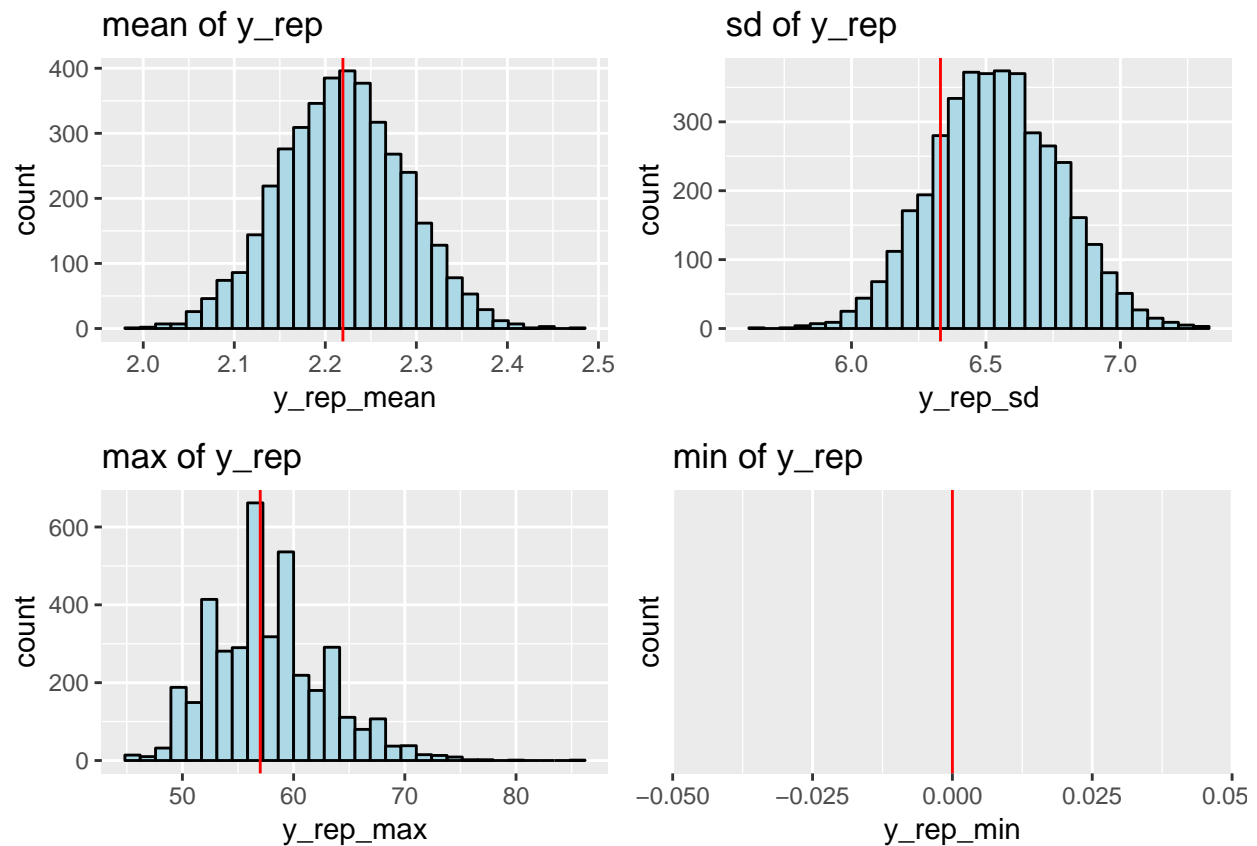
geom_histogram(fill='lightblue', color='black') +
geom_vline(xintercept = sd(y), color='red') +
ggtitle('sd of y_rep')

df <- data.frame(y_rep_max = apply(X = sims$y_rep, MARGIN = 1, FUN = max))
maxgg <- ggplot(df, aes(x=y_rep_max)) +
  geom_histogram(fill='lightblue',color='black') +
  geom_vline(xintercept = max(y), color='red') +
  ggtitle('max of y_rep')

df <- data.frame(y_rep_min = apply(X = sims$y_rep, MARGIN = 1, FUN = min))
mingg <- ggplot(df, aes(x=y_rep_min)) +
  geom_histogram(fill='lightblue',color='black') +
  geom_vline(xintercept = min(y), color='red') +
  ggtitle('min of y_rep')

gridExtra::grid.arrange(meangg, sdgg, maxgg, mingg,
  layout_matrix = rbind(c(1, 2),
    c(3, 4)))


```



```

# Scatterplot of two test statistics
ppc_scatter_avg(y = y, yrep = y_rep)

```



Evaluation (RMSE)

```
## K fold CV
set.seed(1234)
splited_inputs <- split(inputs, sample(rep(1:5, 176)))

for (i in 1:5){
  a <- c(1,2,3,4,5)[-i]
  inputs_test = splited_inputs[[i]]
  inputs_train = rbind(splited_inputs[[a[1]]],
                       splited_inputs[[a[2]]],
                       splited_inputs[[a[3]]],
                       splited_inputs[[a[4]]])

  y_train = inputs_train$sum_y
  y_test = inputs_test$sum_y

  ## Inputs for STAN
  n_city_train = inputs_train$n_city
  n_city_test = inputs_test$n_city

  X_train = inputs_train %>%
    dplyr::select(-ID_state, -n_city, -sum_y, -sum_y2)
  X_test = inputs_test %>%
    dplyr::select(-ID_state, -n_city, -sum_y, -sum_y2)
```

```

N_train = nrow(X_train)
N_test = nrow(X_test)

D = ncol(X_train)

baseline_data = list(N_train=N_train, N_test=N_test, D=D,
                     X_train=X_train, X_test=X_test,
                     n_city_train = n_city_train,
                     n_city_test = n_city_test,
                     y_train = y_train)
fit_cv <- sampling(baseline_model_binom, data=baseline_data, seed=1234)
name = paste('modell1_cv', as.character(i), '.rds', sep = "")
saveRDS(fit_cv, file = name)
}

```

```

rmse <- c()
for (i in 1:5){
  name = paste('modell1_cv', as.character(i), '.rds', sep = "")
  fit_cv <- readRDS(file = name)
  sims_cv <- rstan::extract(fit_cv)
  y_hat <- apply(X = sims_cv$y_rep_cv, MARGIN = 2, FUN = median)

  test_df = data.frame(ID_state = inputs_test$ID_state,
                      y_test = y_test,
                      y_hat = y_hat)

  test_df <- test_df %>%
    summarize(y_sum_test = sum(y_test),
              y_sum_hat = sum(y_hat)) %>%
    arrange(desc(y_sum_test)) %>%
    ungroup()

  #mse_baseline = mean((test_df$y_sum_test) ** 2)
  rmse[i] = sqrt(mean((test_df$y_sum_hat - test_df$y_sum_test) ** 2))
}

```

```

average_RMSE <- mean(rmse)
sd_RMSE <- sd(rmse)

cat('Model RMSE average: ', average_RMSE)

```

Model RMSE average: 75

```
cat('Model RMSE standard deviation: ', sd_RMSE)
```

Model RMSE standard deviation: 33